

UC Berkeley  
Teaching Professor  
Dan Garcia

# CS61C

## Great Ideas in Computer Architecture (a.k.a. Machine Structures)

## C Pointers, Arrays, and Strings

# More C Features

- More C Features
- Pointers
- Using Pointers Effectively
- Arrays, Pointer Arithmetic
- Array Pitfalls
- Strings
- Endianness, Word Alignment

# The C bool: True or False?

- Boolean is **not** a primitive type in C! Instead:
- FALSE:
  - **0** (integer, i.e., all bits are 0)
  - **NULL** (pointer) (more later)
- TRUE:
  - **Everything else!**
  - Note: Same is true in Python.
- Nowadays: true and false are provided by **stdbool.h**.

```
if (42) {  
    printf("meaning of life\n");  
}
```

- A. meaning of life
- B. (nothing)

# More Typing: Typedefs and Structs

- `typedef` allows you to define new types.

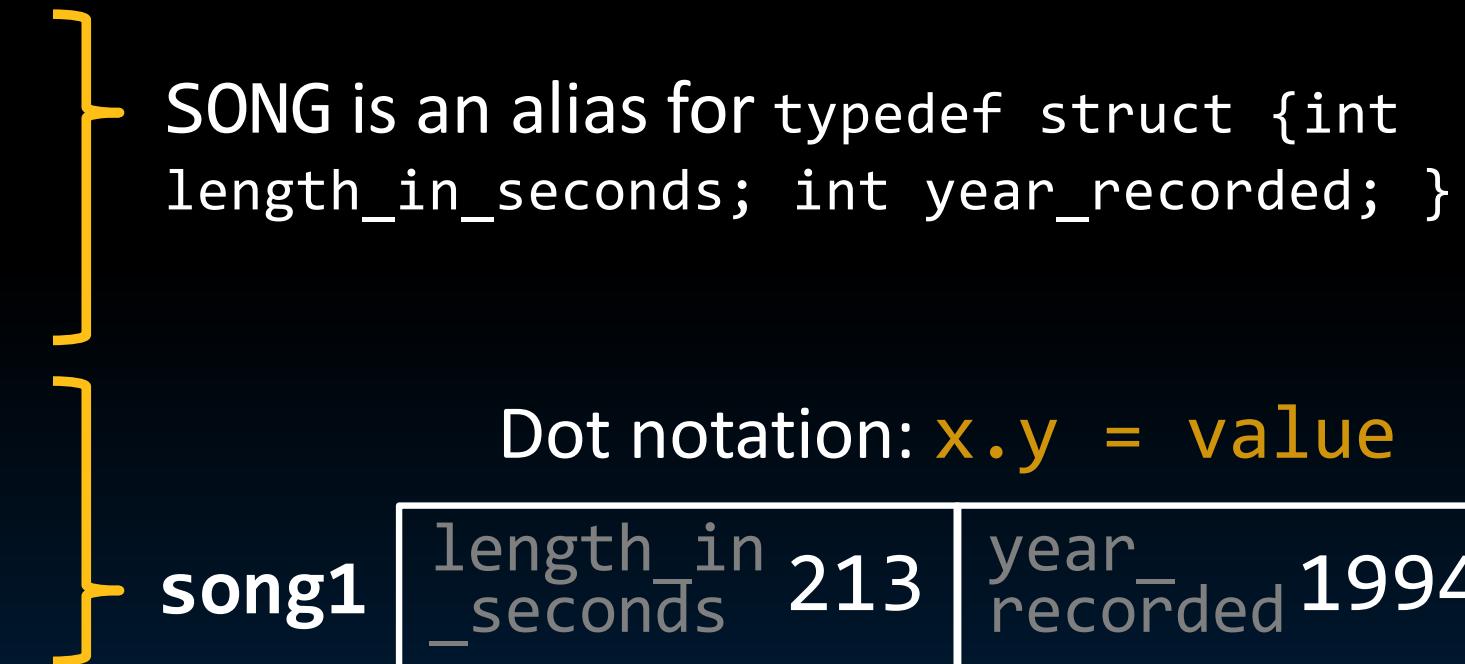
```
typedef uint8_t BYTE;  
BYTE b1, b2;
```

- structs are structured groups of variables, e.g.,

```
typedef struct {  
    int length_in_seconds;  
    int year_recorded;  
} SONG;
```

```
SONG song1;  
song1.length_in_seconds = 213;  
song1.year_recorded = 1994;
```

```
SONG song2;  
song2.length_in_seconds = 248;  
song2.year_recorded = 1988;
```



Structs are **not** objects!  
The `.` operator is **not** a  
method call! (more later)

- Constant, **const**, is assigned a typed value once in the declaration.

- Value can't change during entire execution of program.

```
const float golden_ratio = 1.618;
const int    days_in_week = 7;
const double the_law      = 2.99792458e8;
```

- You can have a constant version of any of the standard C variable types.
- #define PI (3.14159) is a CPP (C Preprocessor) Macro.
  - Prior to compilation, preprocess by performing string replacement in the program based on all #define macros.
  - Replace all PI with (3.14159) → In effect, makes PI a “constant”
- Enums: a group of related integer constants. E.g.,

```
enum cardsuit {CLUBS, DIAMONDS, HEARTS, SPADES};
enum color {RED, GREEN, BLUE};
```

# Pointers

- More C Features
- Pointers
- Using Pointers Effectively
- Arrays, Pointer Arithmetic
- Array Pitfalls
- Strings
- Endianness, Word Alignment

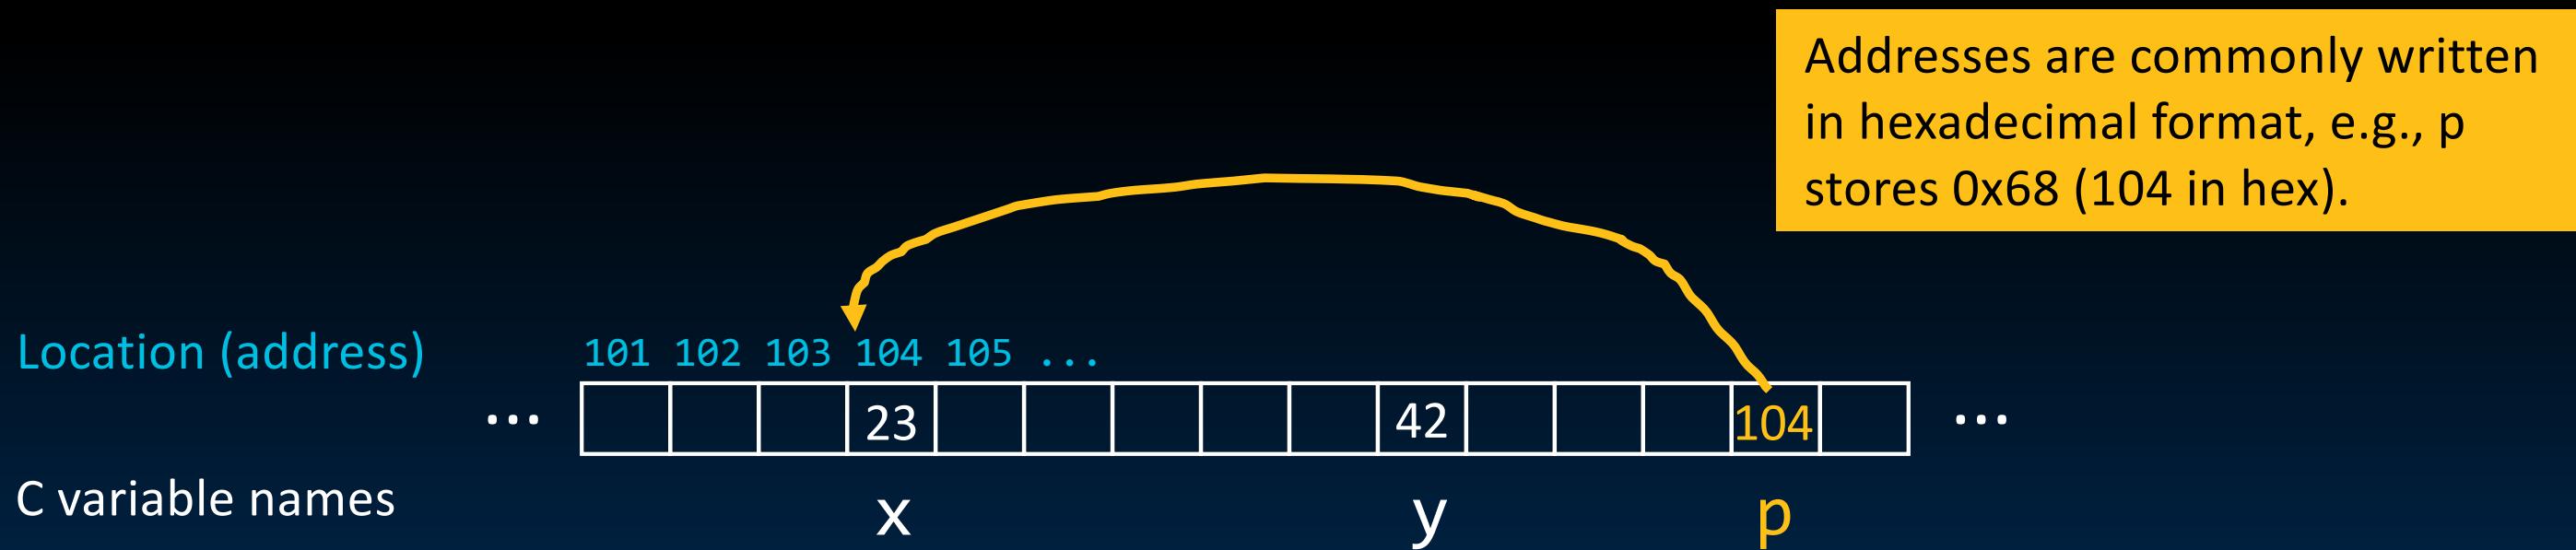
# Memory Is a Single Huge Array

- Consider memory to be a byte-addressed array.
  - Each cell of the array has an address associated with it.
  - Each cell also stores some value.
- Don't confuse the **address** referring to a memory location with the **value** stored in that location.
- For now, the abstraction lets us think we have access to  $\infty$  memory, numbered from 0...



# Pointers store addresses

- An **address** refers to a particular memory location.
- In other words, it “points” to a memory location.
- **Pointer:** A variable that contains the address of another variable.



# Pointer Syntax

0x100

p

???

0x104

x

3

```
1 int *p;  
2 int x = 3;
```

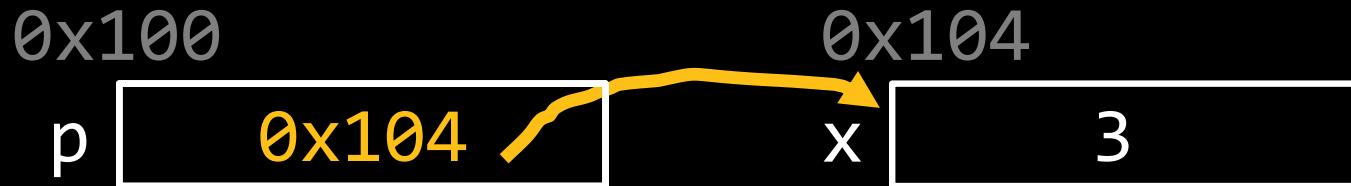
```
3 p = &x;
```

```
4 printf("p points to %d\n",  
       *p);
```

```
5 *p = 5;
```

- Declaration
- Tells compiler that **variable p is address** of an int

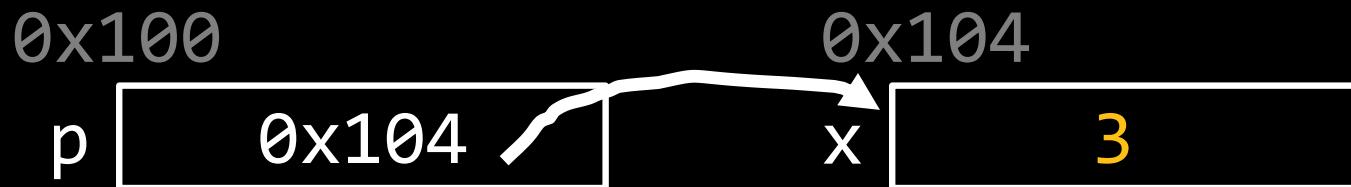
# Pointer Syntax



```
1 int *p;  
2 int x = 3;  
3 p = &x;  
4 printf("p points to %d\n",  
       *p);  
5 *p = 5;
```

- Declaration
- Tells compiler that **variable p** is **address of an int**
- Tells compiler to assign **address of x** to p
- &: “**address operator**” in this context

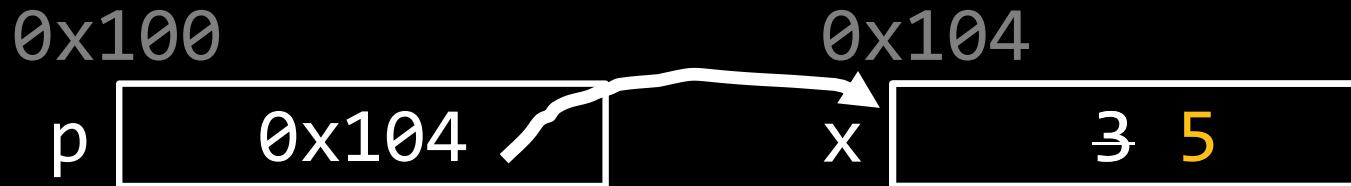
# Pointer Syntax



```
1 int *p;
2 int x = 3;
3 p = &x;
4 printf("p points to %d\n",
      *p);
5 *p = 5;
```

- Declaration
- Tells compiler that **variable p is address of an int**
- Tells compiler to assign **address of x** to p
- &: “**address operator**” in this context
- Gets **value pointed to by p**
- \*: “**dereference operator**” in this context

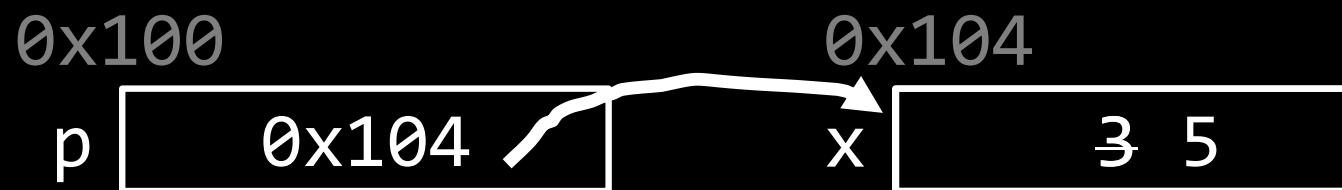
# Pointer Syntax



```
1 int *p;
2 int x = 3;
3 p = &x;
4 printf("p points to %d\n",
      *p);
5 *p = 5;
```

- Declaration
- Tells compiler that **variable p is address of an int**
- Tells compiler to assign **address of x** to p
- &: “**address operator**” in this context
- Gets value pointed to by p
- \*: “**dereference operator**” in this context
- Changes value pointed to by p
- Use deref operator \* on left of =

# Pointer Syntax



```
1 int *p;
2 int x = 3;
3 p = &x;
4 printf("p points to %d\n",
      *p);
5 *p = 5;
```

The “\*” is used in two ways:

**Declaration (L1):** Indicate p is a pointer

**Dereference (L4,L5):** Value pointed to by p

- Declaration
- Tells compiler that **variable p is address of an int**
  
- Tells compiler to assign **address of x** to p
- &: “**address operator**” in this context
  
- Gets **value pointed to by p**
- \*: “**dereference operator**” in this context
  
- Changes **value pointed to by p**
- Use deref operator \* on left of =

# Pointers are Useful When Passing Parameters

- C passes parameters “**by value**”:
  - A function parameter gets assigned a copy of the argument value.

**Changing the function's copy  
cannot change the original.**

```
void addOne (int x)
{
    x = x + 1;
}
```



```
int y = 3;
addOne(y);
```



⚠ **y is still 3...**

# Pointers are Useful When Passing Parameters

- C passes parameters “by value”:
  - A function parameter gets assigned a copy of the argument value.

Changing the function's copy cannot change the original.

```
void addOne (int x)
{
    x = x + 1;
}

int y = 3;
addOne(y);
```



⚠️ y is still 3...

To get a function to change a value, **pass in a pointer**.

```
void addOne (int *p)
{
    *p = *p + 1;
}

int y = 3;
addOne(&y);
```



✓ y is now 4!

Garcia, Kao

# Pointers in C ... The Good, Bad, and the Ugly

## Why use pointers?

- To pass a large struct or array to a function, it's easier/faster/etc. to pass a pointer.
  - Otherwise, we'd need to copy a huge amount of data!
- At the time C was invented (early 1970s), compilers didn't produce efficient code, so C was designed to give human programmer more flexibility.
  - Nowadays, computers are 100,000x faster; compilers are also way, way, way better.
- Still used for low-level system code, as well as implementation of "pass-by-reference" object paradigms in other languages.
- In general, pointers allow cleaner, more compact code.

## ⚠ So, what are the drawbacks?

- Pointers are probably the single largest source of bugs in C. **Be careful!**
  - Most problematic with dynamic memory management
  - Dangling references and memory leaks

} (more later)

# Common C Bug: Garbage Addresses



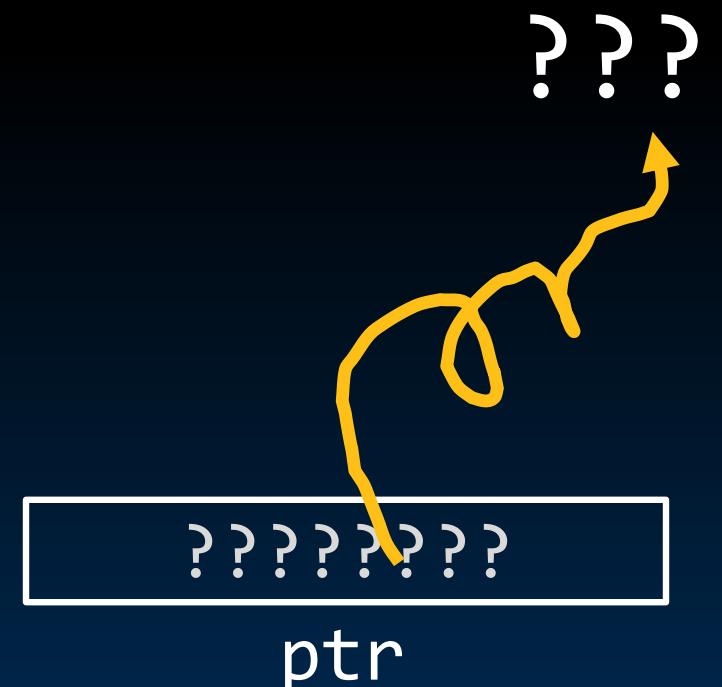
- Declaring a pointer just allocates space to hold the pointer.
  - It does not allocate something to be pointed to!

- Recall: **Local variables in C are not initialized.**

They may contain *anything*.

- What does the following code do?

```
void f()
{
    int *ptr;
    *ptr = 5;
}
```



# Using Pointers Effectively

- More C Features
- Pointers
- Using Pointers Effectively
- Arrays, Pointer Arithmetic
- Array Pitfalls
- Strings
- Endianness, Word Alignment

# Pointers to Different Data Types

- Pointers are used to point to a variable of a particular data type.
  - Normally a pointer can only point to one type.
- `void *` is a type that can point to anything (generic pointer).
  - Use sparingly to help avoid program bugs... and security issues... and a lot of other bad things!
- You can even have pointers to functions...
  - `int (*fn) (void *, void *) = &foo;`
    - `fn` is a function that accepts two `void *` pointers and returns an `int` and is initially pointing to the function `foo`.
    - `(*fn)(x, y);` will then call the function

```
int *xptr;  
char *str;  
struct llist *foo_ptr;
```

(more later)



# NULL pointers...

- The pointer of all 0s is special.
  - The "NULL" pointer, like in Java, Python, etc...
- If you write to or read from a null pointer, your program should crash.
- Since "0 is false", its very easy to do tests for null:
  - `if(!p) { /* p is a null pointer */ }`
  - `if(q) { /* q is not a null pointer */ }`

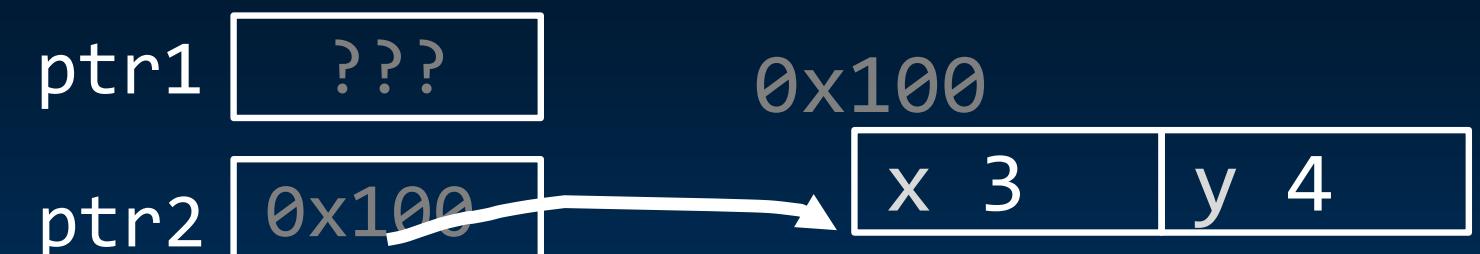
```
char *p = NULL;
```

```
p 0x00000000
```

# Structs, Revisited

```
typedef struct {  
    int x;  
    int y;  
} Coord;  
  
/* declarations */  
Coord coord1, coord2;  
Coord *ptr1, *ptr2;  
  
/* instantiations  
go here... */
```

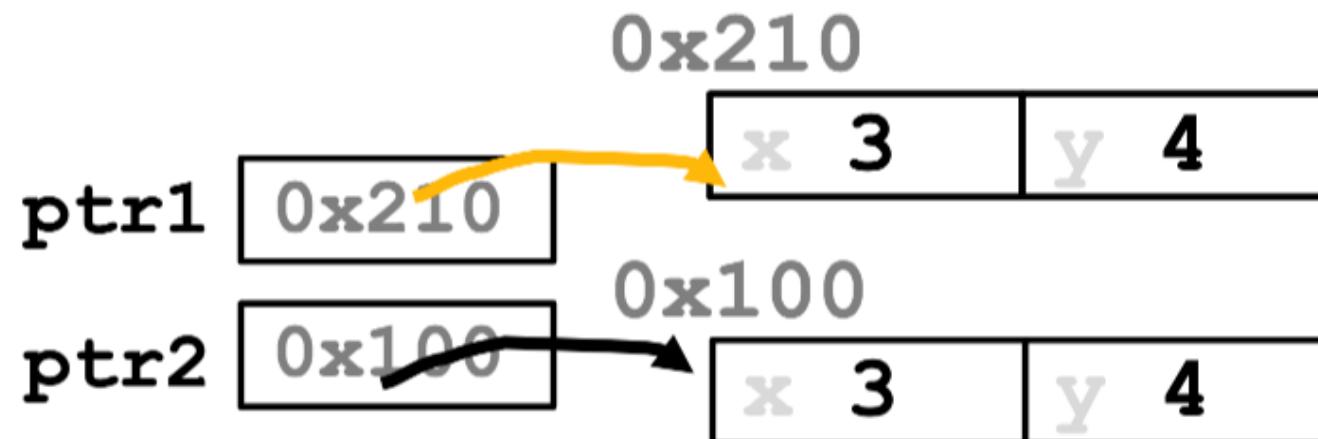
```
/* dot notation */  
int h = coord1.x;  
coord2.y = coord1.y;  
  
/* arrow notation = deref + struct access */  
int k;  
k = (*ptr1).x;  
k = ptr1->x; // equivalent  
  
/* This compiles, but what does it do? */  
ptr1 = ptr2;
```



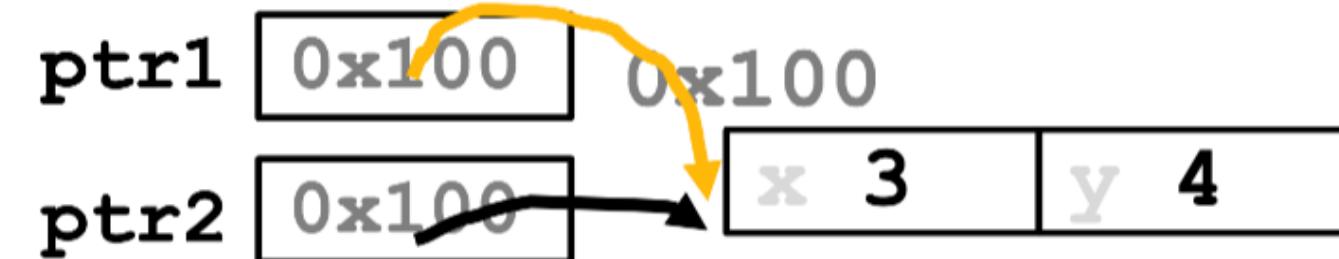
Assume ptr2 points to an initialized Coord struct {3, 4}.  
What does `ptr1 = ptr2;` do?

0

CHOICE A  
(click on this side)



CHOICE B  
(click on this side)



# Pointers and Structures

```
typedef struct {  
    int x;  
    int y;  
} Coord;  
  
/* declarations */  
Coord coord1, coord2;  
Coord *ptr1, *ptr2;  
  
/* instantiations  
go here... */
```

```
/* dot notation */  
int h = coord1.x;  
coord2.y = coord1.y;
```

```
/* arrow notation */  
int k;  
k = ptr1->x;  
k = (*ptr1).x; // equivalent
```

/\* This compiles, but what does it do? \*/  
ptr1 = ptr2;



# Arrays, Pointer Arithmetic

- More C Features
- Pointers
- Using Pointers Effectively
- Arrays, Pointer Arithmetic
- Array Pitfalls
- Strings
- Endianness, Word Alignment

# A C array is really just a big block of memory

- Declaration:

- `int arr[2];`
- ...declares a 2-element integer array

???	???
-----	-----

arr

- Declaration and initialization

- `int arr[] = {795, 635};`
- declares and fills a 2-elt integer array

795	635
-----	-----

arr

- Accessing elements:

- `arr[num]`
- returns the num<sup>th</sup> element.
- This is shorthand for **pointer arithmetic**.

`arr[0]; // 795`

# Pointer Arithmetic

Equivalent:

$$a[i] \Leftrightarrow *(a+i)$$

pointer + n

Adds  $n * \text{sizeof}(\text{"whatever pointer is pointing to"})$  to memory address.

pointer - n

Subtracts  $n * \text{sizeof}(\text{"whatever pointer is pointing to"})$  from memory address.

`*sizeof`: compile-time op for # of bytes in object.



```
// 32-bit unsigned int array
uint32_t arr[] = {50, 60, 70};

uint32_t *q = arr;
```

# Pointer Arithmetic

Equivalent:

$$a[i] \Leftrightarrow *(a+i)$$

pointer + n

Adds  $n * \text{sizeof}(\text{"whatever pointer is pointing to"})$  to memory address.

pointer - n

Subtracts  $n * \text{sizeof}(\text{"whatever pointer is pointing to"})$  from memory address.

`*sizeof`: compile-time op for # of bytes in object.



```
// 32-bit unsigned int array
uint32_t arr[] = {50, 60, 70};

uint32_t *q = arr;
```

```
printf("    *q: %d is %d\n", *q, q[0]);
printf("*q+1: %d is %d\n", *(q+1), q[1]);
printf("*q-1: %d is %d\n", *(q-1), q[-1]);
```

```
*q: 50 is 50
*q+1: 60 is 60
*q-1: ??? is ???
```

How to get a function to change a pointer?

- Suppose we want `increment_ptr` to change where `q` points to.

Remember: C is pass-by-value!

What gets printed?

\*`q` is 50

```
1 void increment_ptr(int32_t *p)
2 {   p = p + 1;   }
3 int32_t arr[3] = {50, 60, 70};
4 int32_t *q = arr;
5 increment_ptr(q);
6 printf("*q is %d\n", *q);
```

0x100  
0x104

p

0x100 0x104 0x108 0x10c ... 0x120





How to get a function to change a pointer?

- Suppose we want `increment_ptr` to change where `q` points to.

Remember: C is pass-by-value!

- Instead, pass a **pointer to a pointer** ("handle").
- Declared as `data_type **h`.

Now, what gets printed?

\*`q` is 60

```
1 void increment_ptr(int32_t **h)
2 {   *h = *h + 1;   }
3
4 int32_t arr[3] = {50, 60, 70};
5 int32_t *q = arr;
6 increment_ptr(&q);
7 printf("*q is %d\n", *q);
```

0x120  
h



# Array Pitfalls

- More C Features
- Pointers
- Using Pointers Effectively
- Arrays, Pointer Arithmetic
- Array Pitfalls
- Strings
- Endianness, Word Alignment



- Declare array and initialize all elements of an array of known size **n**:
  - **⚠ Wrong**

```
int i, ar[10];
for(i = 0; i < 10; i++){ ... }
```
  - **Strongly encouraged**

```
int ARRAY_SIZE = 10;
int i, a[ARRAY_SIZE];
for(i = 0; i < ARRAY_SIZE; i++){ ... }
```
- **Why? SINGLE SOURCE OF TRUTH!**
  - Utilize indirection and avoid maintaining two copies of the number 10!

# Arrays vs Pointers

- Arrays are (almost) identical to pointers.
  - `char *string` and `char string[]` are nearly identical declarations
  - They differ in **very subtle** ways: incrementing, declaration of filled arrays...(more in a bit)
- Accessing Array Elements
  - `arr` is an array variable, but it looks like a pointer in many respects (though not all).
  - `arr[0]` is the same as `*arr`
  - `arr[2]` is the same as `*(arr+2)`

An array variable is a “pointer” to the first (0-th) element.

# Arrays are not implemented as you'd think...

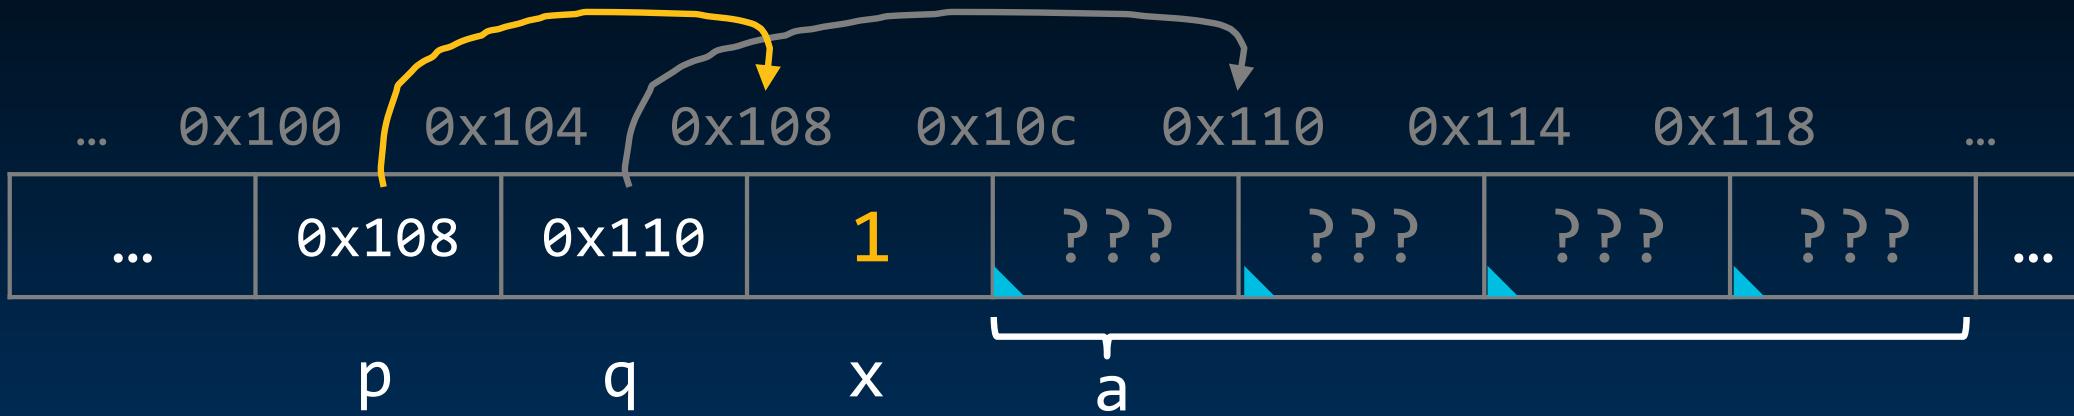
```
int *p, *q, x;  
int a[4];  
p = &x;  
q = a + 1;
```

```
*p:1, p:108, &p:100
```

```
{  
    *p = 1;  
    printf("*p:%d, p:%x, &p:%x\n", *p, p, &p); // %d: signed decimal, %x: hex
```

```
*q = 2;  
printf("*q:%d, q:%x, &q:%x\n", *q, q, &q);
```

```
*a = 3;  
printf("*a:%d, a:%x, &a:%x\n", *a, a, &a);
```



# Arrays are not implemented as you'd think...

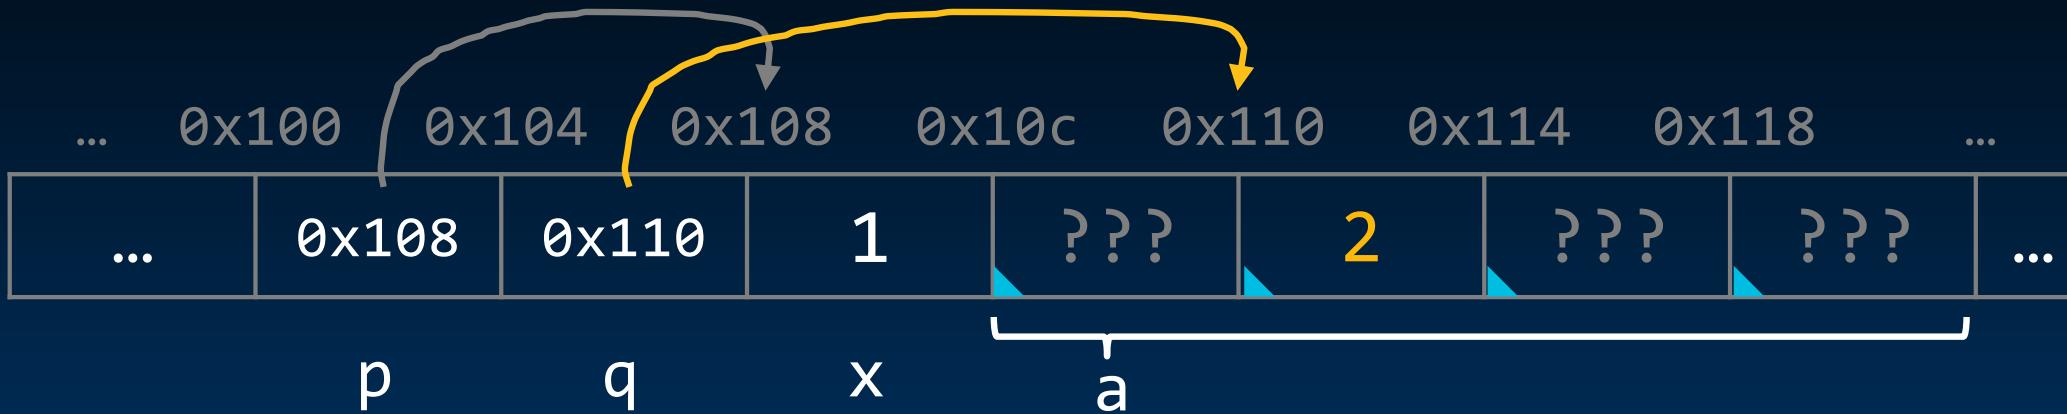
```
int *p, *q, x;  
int a[4];  
p = &x;  
q = a + 1;
```

```
*p:1, p:108, &p:100  
*q:2, q:110, &q:104
```

```
*p = 1;  
printf("*p:%d, p:%x, &p:%x\n", *p, p, &p); // %d:signed decimal,%x:hex
```

```
{  
*q = 2;  
printf("*q:%d, q:%x, &q:%x\n", *q, q, &q);
```

```
*a = 3;  
printf("*a:%d, a:%x, &a:%x\n", *a, a, &a);
```



# Arrays are not implemented as you'd think...

```
int *p, *q, x;  
int a[4];  
p = &x;  
q = a + 1;
```

```
*p:1, p:108, &p:100  
*q:2, q:110, &q:104  
*a:3, a:10c, &a:10c
```

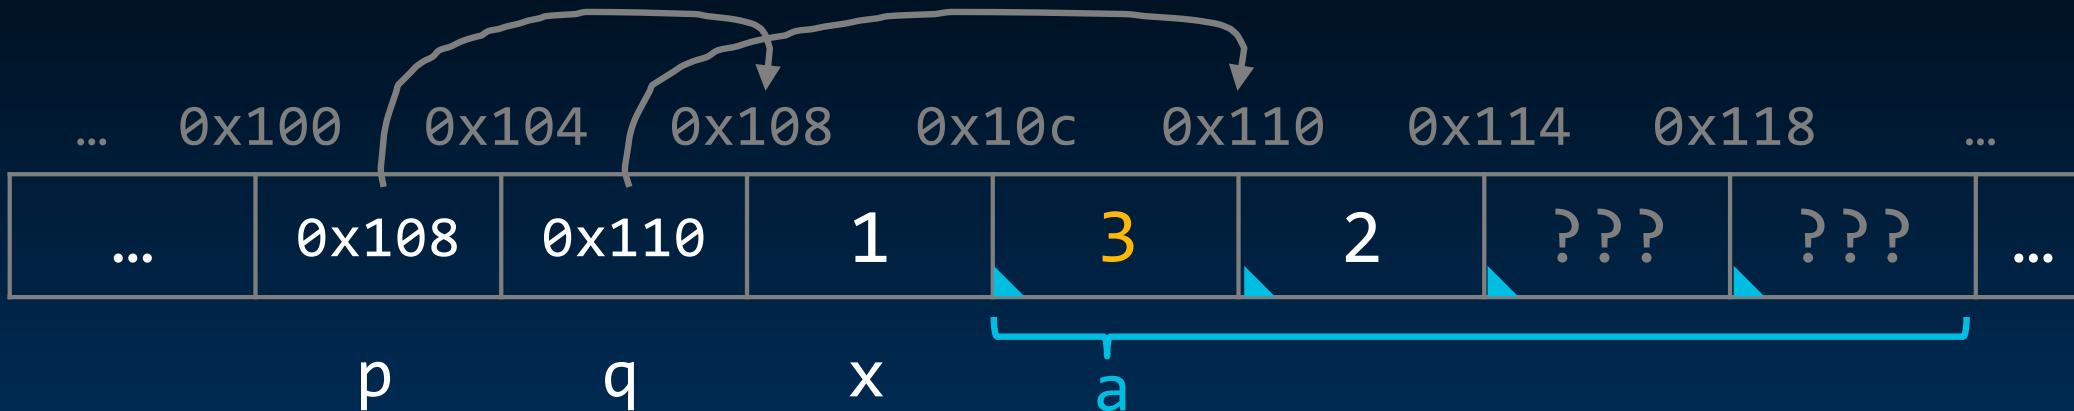


```
*p = 1;  
printf("*p:%d, p:%x, &p:%x\n", *p, p, &p); // %d:signed decimal,%x:hex
```

```
*q = 2;  
printf("*q:%d, q:%x, &q:%x\n", *q, q, &q);
```

```
{ *a = 3;  
printf("*a:%d, a:%x, &a:%x\n", *a, a, &a);
```

K&R: “An array name is not a variable”



Dan: “A C array is really just a big block of memory”

# Array Are Very Primitive (1/3)



1. Array bounds are not checked during element access.
  - Consequence: We can accidentally access off the end of an array!

```
int ARRAY_SIZE = 100;
int foo[ARRAY_SIZE];
int i;
.....
for(i = 0; i <= ARRAY_SIZE; ++i) {
    foo[i] = 0;
}
```

- Corrupts other parts of the program...
  - Including internal C data
- May cause crashes later.



The image shows a screenshot of a WhatsApp message. The header 'WhatsApp Security Advisories' is in bold black font. Below it, 'CVE-2019-11933' is listed. A blue box highlights the text: 'A heap buffer overflow bug in libpl\_droidsonroids\_gif before 1.2.19, as used in WhatsApp for Android before version 2.19.291 could allow remote attackers to execute arbitrary code or cause a denial of service.' At the bottom, there is a note: 'click a link preview from a specially crafted text message.'

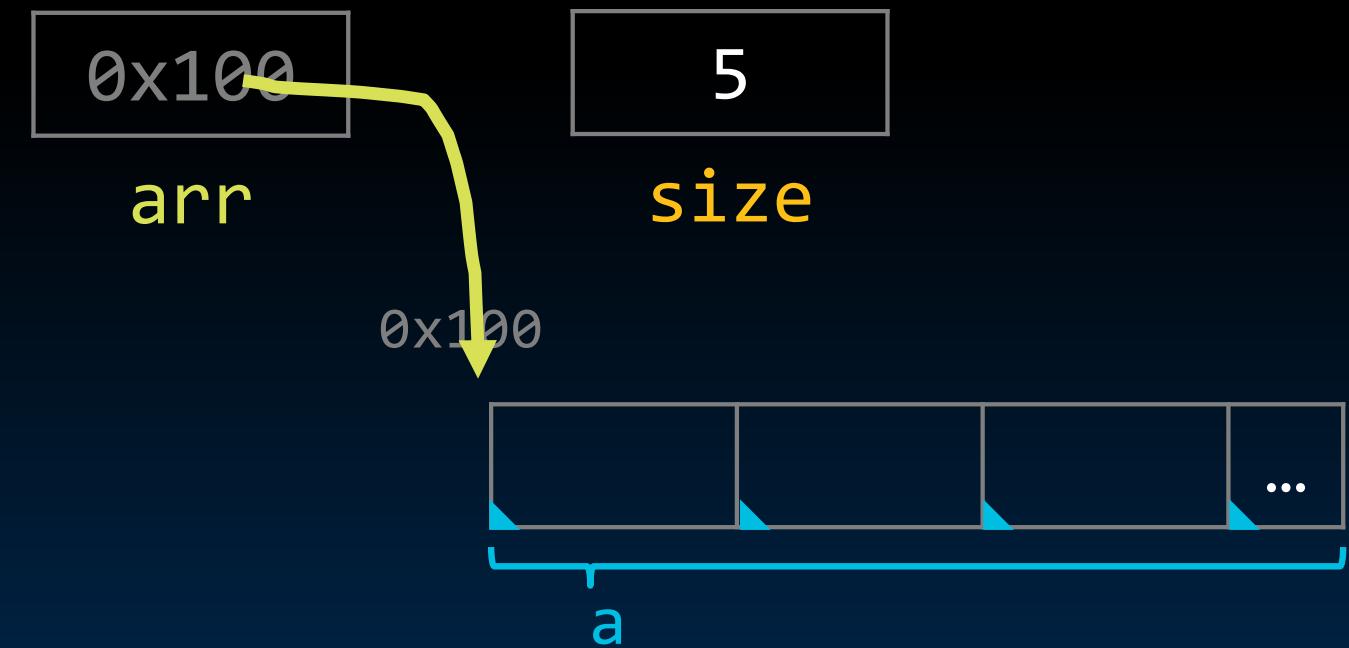
# Array Are Very Primitive (2/3)



1. Array bounds are not checked during element access.
  - Consequence: We can accidentally access off the end of an array!
2. An array is passed to a function as a pointer.
  - Consequence: The array size is lost! Be careful with `sizeof()`!

same as: `int *arr`

```
int bar(int arr[], unsigned int size)
{
    ... arr[size - 1] ...
}
int main(void)
{
    int a[5], b[10];
    ...
    bar(a, 5);
    ...
}
```



- You should always explicitly **include array length as a parameter**.

# Array Are Very Primitive (3/3)



1. Array bounds are not checked during element access.
  - Consequence: We can accidentally access off the end of an array!
2. An array is passed to a function as a pointer.
  - Consequence: The array size is lost! Be careful with `sizeof()`!
3. Declared arrays are only allocated while the scope is valid.

```
char *foo() {  
    char string[32]; ...;  
    return string;  
} is incorrect
```

Solution:  
Dynamic memory allocation!  
[\(more later\)](#)

# Array Are Very Primitive, Summary

1. Array bounds are not checked during element access.
  - Consequence: We can accidentally access off the end of an array!
2. An array is passed to a function as a pointer.
  - Consequence: The array size is lost! Be careful with `sizeof()`!
3. Declared arrays are only allocated while the scope is valid.

## Segmentation faults and bus errors:

- These are VERY difficult to find; be careful!
- You'll learn how to debug these in Lab 02 with `gdb`...

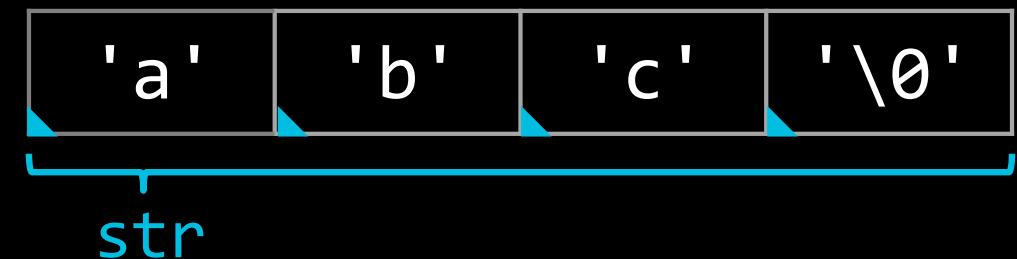
# Strings

- More C Features
- Pointers
- Using Pointers Effectively
- Arrays, Pointer Arithmetic
- Array Pitfalls
- Strings
- Endianness, Word Alignment

- A C string is just an array of characters, followed by **a null terminator**.
  - Null terminator: the byte 0 (number) aka '\0' character)
- The standard C library **string.h** assumes **null-terminated strings**.
  - String operations do **not** include the null terminator, e.g., when you ask for length of a string!

Lab 01: string.h: strcpy() vs strncpy()

```
char str[] = "abc";
```



... `strlen(str)` ... // 3 !

```
// possible implementation*
int strlen(char s[])
{
    int n = 0;
    while (*(s++) != 0) { n++; }
    return n;
}
```

\*for actual `strlen()` implementation, see [glibc](#)

# Endianness, Word Alignment

- More C Features
- Pointers
- Using Pointers Effectively
- Pointer Arithmetic
- Array Pitfalls
- Strings
- Endianness, Word Alignment

# Memory and Addresses

How to read byte addresses:

0x...FFFE

0x...FFFC

- Modern machines are “byte-addressable.”
  - Hardware’s memory composed of 8-bit storage cells; each byte has a unique address
- We commonly think in terms of “word size”:
  - aka number of bits in an address
  - A **32b** architecture has **4-byte words**.
  - All pointer sizes on 32b architecture:  
`sizeof(int *) == ...  
== sizeof(char *) == 4`

	+3	+2	+1	+0
0xFFFFFFF4				
0xFFFFFFF8				
0xFFFFFFFF0	xx	xx	xx	xx
0xFFFFFFFF4				
0xFFFFFFFF8				
0xFFFFFFFFC	int32_t *			
0xFFFFFFFFE0	xx	xx	xx	xx
0xFFFFFFFFE4	short *			
0xFFFFFFFFE8	char *			
...	xx	xx	xx	xx
0x00	xx	xx	xx	xx
0x04	xx	xx	xx	xx
0x08	xx	xx	xx	xx
0x0C	xx	xx	xx	xx

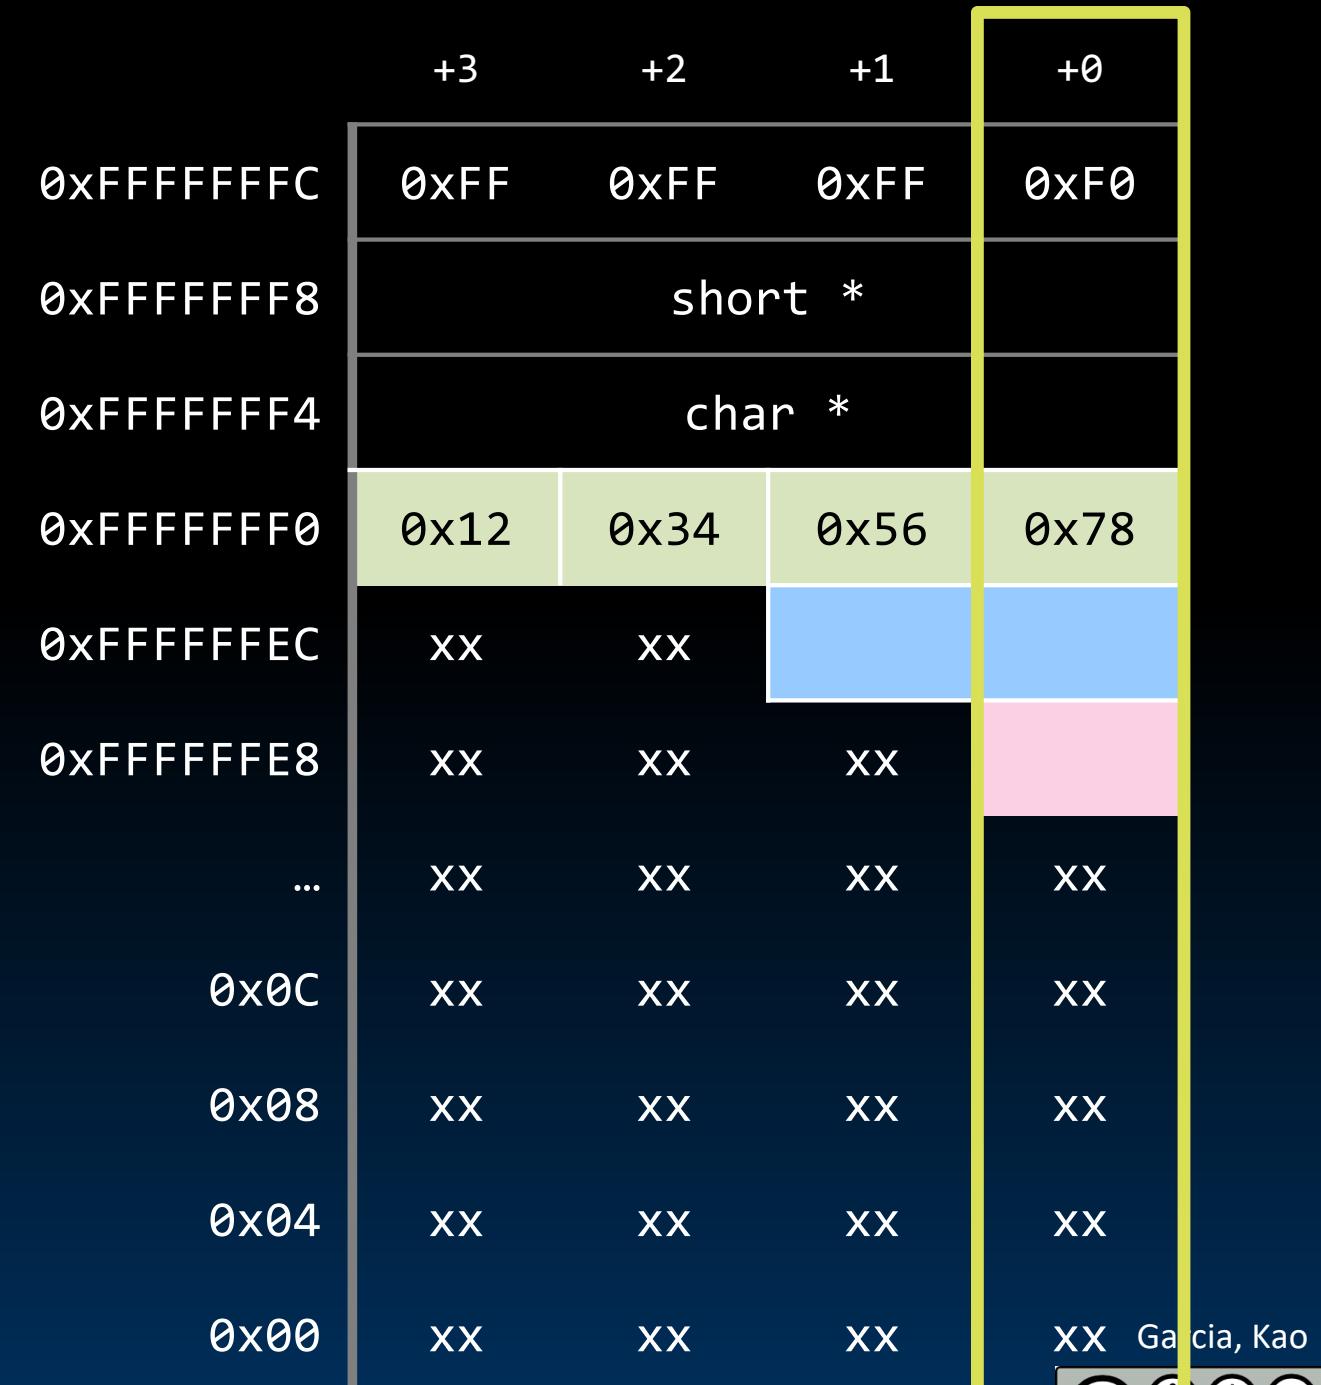


# Endianness

- The hive machines are “little endian”.
  - The **least significant byte** of a value is stored first.
- (Contrast with “big endian”)
  - (Most significant byte is stored first)

**int32\_t\*** pointer  
0xFFFFFFFF0

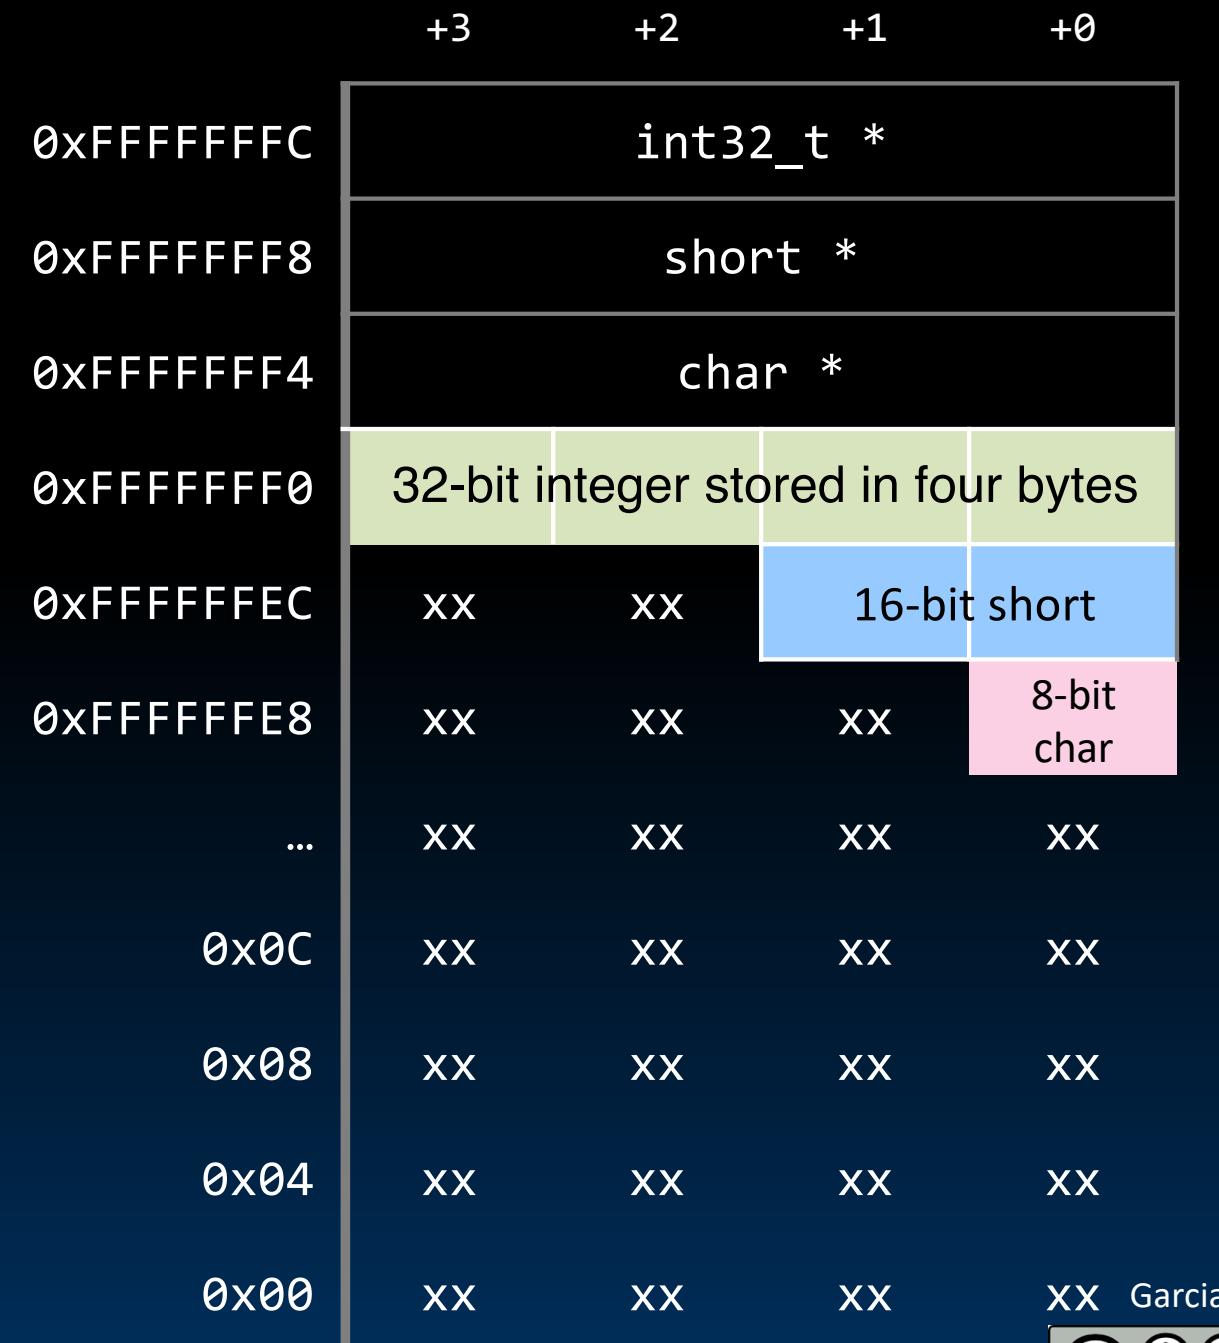
**int32\_t** value  
0x12345678



# Word Alignment

Read for HW01

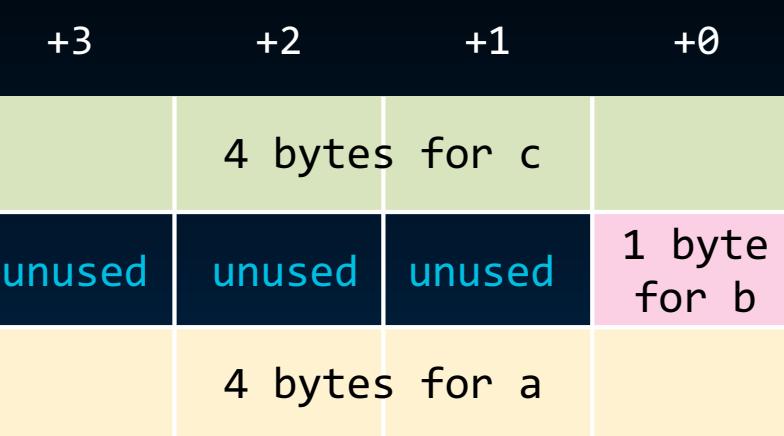
- We often want “word alignment”:
  - Some processors will not allow you to address **32b** values without being on **4-byte** boundaries.
  - Others will just be very slow if you try to access “unaligned” memory.



- A “struct” is really an instruction to C on how to arrange a bunch of bytes in a bucket.

```
struct foo {  
    int32_t a;  
    char b;  
    struct foo *c;  
}
```

- Structs provide enough space for the data.
- C compilers often **align** the data with **padding**.
- For this struct, the actual layout on a 32b architecture would be as follows.
  - Note the 3 bytes of padding
  - `sizeof(struct foo) == 12`



# And in Conclusion...

- C pointers and arrays are **pretty much the same**, except with function calls.
- C knows how to **increment pointers**.
- C is an efficient language, but with little protection.
  - Array bounds **not checked**
  - Variables **not automatically initialized**
- Use handles to change pointers
- (Beware) The cost of efficiency is more overhead for the programmer.
  - “C gives you a lot of extra rope, don’t hang yourself with it!”

