


UC Berkeley
Teaching Professor
Dan Garcia

CS61C


Great Ideas
in
Computer Architecture
(a.k.a. Machine Structures)

Memory (Mis)Management


From ENIAC (1946) to EDSAC (1949)



- ENIAC: First Electronic General-Purpose Computer
- Needed 2-3 days to setup new program
- Programmed with patch cords and switches
 - At that time & before, "computer" mostly referred to people who did calculations
 - Mostly women! (See *Hidden Figures*, 2016)



- EDSAC: First General Stored-Program Computer
- Programs held as **numbers in memory**
 - Revolution! Program is also data!
- 35-bit binary **two's complement** words



What gets printed?


[Concept Check]

sizeof(): compile-time operator; gives size in bytes (of type or variable).

```
// for this exercise, assume
// shorts are 16b on a 64-bit architecture
void mystery(short arr[], int len) {
    printf("%d ", len);
    printf("%d\n", sizeof(arr));
}

int main() {
    short nums[] = {1, 2, 3, 99, 100};
    printf("%d ", sizeof(nums));
    mystery(nums, sizeof(nums)/sizeof(short));
    return 0;
}
```

- A. 10 5 10
- B. 10 5 8
- C. 80 5 80
- D. 80 5 40
- E. Other



What gets printed?

[Concept Check]

sizeof(): compile-time operator; gives size in bytes (of type or variable).

```
// for this exercise, assume
// shorts are 16b on a 64-bit architecture
void mystery(short arr[], int len) {
    printf("%d ", len);
    printf("%d\n", sizeof(arr));
}
```


- A. 10 5 10
- B. 10 5 8
- C. 80 5 80
- D. 80 5 40
- E. Other

Array has decayed
to a pointer

```
int main() {
    short nums[] = {1, 2, 3, 99, 100};
    printf("%d ", sizeof(nums));
    mystery(nums, sizeof(nums)/sizeof(short));
    return 0;
}
```

In array's declared scope,
total array size.


In array's declared scope,
elements in array.



Agenda


Memory Locations

- Memory Locations
- The Stack
- The Heap
- Linked List Example
- When Memory Goes Bad
- Implementing Memory Management




C Program Address Space

- A program's address space contains 4 regions:
 - **Stack**: local variables inside functions, grows downward
 - **Heap**: space requested via `malloc()`; resizes dynamically, grows upward
 - **Data (Static Data)**: variables declared outside main, does not grow or shrink
 - **Text (aka code)**: program executable loaded when program starts, does not change
 - `0x00000000` chunk is unwriteable/unreadable so you that crash on NULL pointer access
- Programming in C requires knowing where objects are in memory, otherwise things don't work as expected.
 - By contrast, Java hides location of objects.



For now, OS somehow prevents accesses between stack and heap.
(more later w/virtual memory)

Garcia, Kao




Where are variables allocated?


- **Global:** If declared **outside** a function, allocated in **data (static)** storage.
- **Local:** If declared **inside** function, allocated on the **stack** and freed when function returns.
 - NB: `main()` is also a function.
- For both these memory types, the management is **automatic**.
 - You don't need to worry about deallocating when you are no longer using them.
 - But a variable **does not exist anymore** once a function ends!

```
int myGlobal;  
... main() {  
    int myTemp;  
    ...  
}
```

0xFFFF FFFF

0x0000 0000






Agenda

The Stack


- Memory Locations
- The Stack
- The Heap
- Linked List Example
- When Memory Goes Bad
- Implementing Memory Management




The Stack

- Every time a function is called, a new **stack frame** is allocated on the stack.
- A stack frame includes:
 - Return “instruction” address (who called me?)
 - Arguments*
 - Space for other local variables
- Stack frames contiguous blocks of memory; the **stack pointer** indicates the start of stack frames.
- When function ends, stack frame is tossed off the stack; frees memory for future stack frames.
- (more later when we cover details for a RISC-V processor architecture)

```
fooA() { fooB(); }  
fooB() { fooC(); }  
fooC() { ... }
```




*see slide notes in pptx for technical elaboration



The Stack is Last In, First Out (LIFO)

```
int main ()
{ a(0); ...
}
void a (int m)
{ b(1);
}
void b (int n)
{ c(2);
}
void c (int o)
{ d(3);
}
void d (int p)
{
}
```




stack

Stack
grows
down



The stack grows
downward; a()'s
local variables have
lower byte addresses
than main()'s, and so
on.

Stack Pointer →



Recall: Array Are Very Primitive

[From last time]




1. Array bounds are not checked during element access.
 - Consequence: We can accidentally access off the end of an array!
2. An array is passed to a function as a pointer.
 - Consequence: The array size is lost! Be careful with `sizeof()`!
3. Declared arrays are only allocated while the scope is valid.

```
char *foo() {  
    char string[32]; ...;  
    return string;  
} is incorrect
```

Solution:

Dynamic memory allocation!
(more late now)




Passing Pointers into the Stack

It is fine to pass a pointer to stack space further down.

- I.e., pointers to addresses higher in the stack point to data in currently allocated stack frames.


```
void load_buf() { ... };
int main() {
    ...
    char buf[...];
    load_buf(buf, BUFLen);
    ...
}
```




However, it is catastrophically bad to return a pointer to something in the stack!

- Memory will be overwritten when other functions are called!
- So your data would no longer exist...and writes can overwrite key pointers, causing crashes!

```
char *make_buf() {
    char buf[50];
    return buf;
}
void foo() {...}
int main() {
    char *stackAddr = \
        make_buf();
    foo(stackAddr);
    ...
}
```






Agenda


The Heap

- Memory Locations
- The Stack
- The Heap
- Linked List Example
- When Memory Goes Bad
- Implementing Memory Management



What is the Heap?

- The heap is dynamic memory – memory that can be allocated, resized, and freed during program runtime.
 - Useful for persistent memory across function calls.
 - But biggest source of pointer bugs, memory leaks, ...
 - Similar to Java **new** command allocates memory....but with key differences below.
- **Huge** pool of mem (usually >> stack), but not allocated in contiguous order.
 - Back-to-back requests for heap memory *could* result in blocks very far apart
- In C, specify number of bytes of memory **explicitly** to allocate/deallocate item.
 - `malloc()`: Allocates raw, uninitialized memory from heap
 - `free()`: Frees memory on heap
 - `realloc()`: Resizes previously allocated heap blocks to new size
 - Unlike the stack, memory gets reused only when programmer **explicitly** cleans up



`void *malloc(size_t n)`

- Allocates a block of uninitialized memory:
 - `size_t n` is an unsigned integer type big enough to “count” memory bytes.
 - Returns `void *` pointer to block of memory on heap.
 - A return of `NULL` indicates no more memory (**always** check for it!!!)

- To allocate a struct:

```
typedef struct { ... } TreeNode;
```

```
TreeNode *tp = (TreeNode *) malloc(sizeof(TreeNode));
```

Typecast casts return value
from type `(void *)` to `(TreeNode *)`

`sizeof(type)` gives size in bytes.


- To allocate an array of 20 ints:

```
int *ptr = (int *) malloc(20*sizeof(int));
```

```
if (ptr != NULL) { ... // always check NULL after
```

- Many years ago ints used to be 16b. Now, 32b or 64b...

Assuming size of objects
can lead to misleading,
unportable code. Use
`sizeof()`!




`void free(void *ptr)`

- Dynamically frees heap memory
 - `ptr` is a pointer containing an address originally returned by `malloc()`/`realloc()`.

```
int *ptr = (int *) malloc (sizeof(int)*20);  
...  
free(ptr); // implicit typecast to (void *)
```

- When you free memory, be sure to pass the original address returned from `malloc()`. Otherwise, crash (or worse!)






`void *realloc(void *ptr, size_t size)`

- Resize a previously allocated block at `ptr` to a new size.
 - Returns new address of the memory block.
 - In doing so, it may need to copy all data to a new location.
 - `realloc(NULL, size);` // behaves like `malloc`
 - `realloc(ptr, 0);` // behaves like `free`, deallocates heap block
- Remember: Always check for return `NULL`, which would mean you've run out of memory!


```
int *ip; ip = (int *) malloc(10*sizeof(int));
... .. /* check for NULL */
ip = (int *) realloc(ip, 20*sizeof(int));
/* contents of first 10 elements retained */
... .. /* check for NULL */
realloc(ip,0); /* equivalent to free(ip); */
```



Agenda

Linked List Example


- Memory Locations
- The Stack
- The Heap
- Linked List Example
- When Memory Goes Bad
- Implementing Memory Management



Linked List Example

```
# include <string.h>
int main() {
1   struct Node *head = NULL;
2   add_to_front(&head, "abc");
   ... // free nodes, strings here...
}
void add_to_front(struct Node **head_ptr, char *data) {
3   struct Node *node = (struct Node *) malloc(sizeof(struct Node));
4   node->data = (char *) malloc(strlen(data) + 1); // extra byte!
5   strcpy(node->data, data); // strcpy also copies null terminator
6   node->next = *head_ptr;
7   *head_ptr = node;
}
```

```
struct Node {
    char *data;
    struct Node *next;
};
```





Linked List Example

```
# include <string.h>
int main() {
1   struct Node *head = NULL;
2   add_to_front(&head, "abc");
   ... // free nodes, strings here...
}

void add_to_front(struct Node **head_ptr, char *data) {
3   struct Node *node = (struct Node *) malloc(sizeof(struct Node));
4   node->data = (char *) malloc(strlen(data) + 1); // extra byte!
5   strcpy(node->data, data); // strcpy also copies null terminator
6   node->next = *head_ptr;
7   *head_ptr = node;
}
```

```
struct Node {
    char *data;
    struct Node *next;
};
```




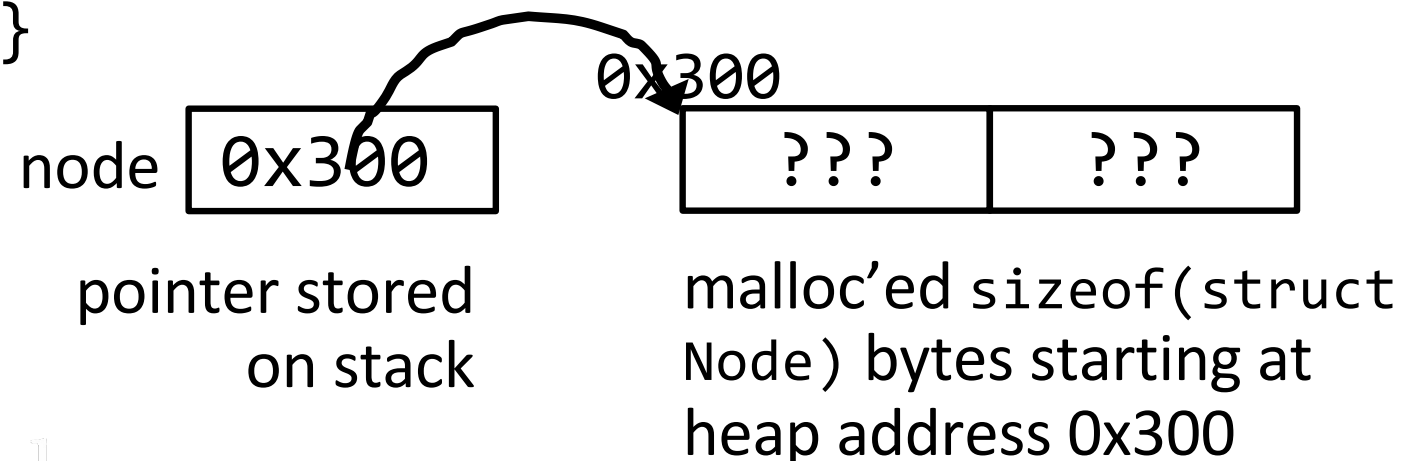



Linked List Example

```
# include <string.h>
int main() {
1   struct Node *head = NULL;
2   add_to_front(&head, "abc");
   ... // free nodes, strings here...
}

void add_to_front(struct Node **head_ptr, char *data) {
3   struct Node *node = (struct Node *) malloc(sizeof(struct Node));
4   node->data = (char *) malloc(strlen(data) + 1); // extra byte!
5   strcpy(node->data, data); // strcpy also copies null terminator
6   node->next = *head_ptr;
7   *head_ptr = node;
}
```

```
struct Node {
    char *data;
    struct Node *next;
};
```



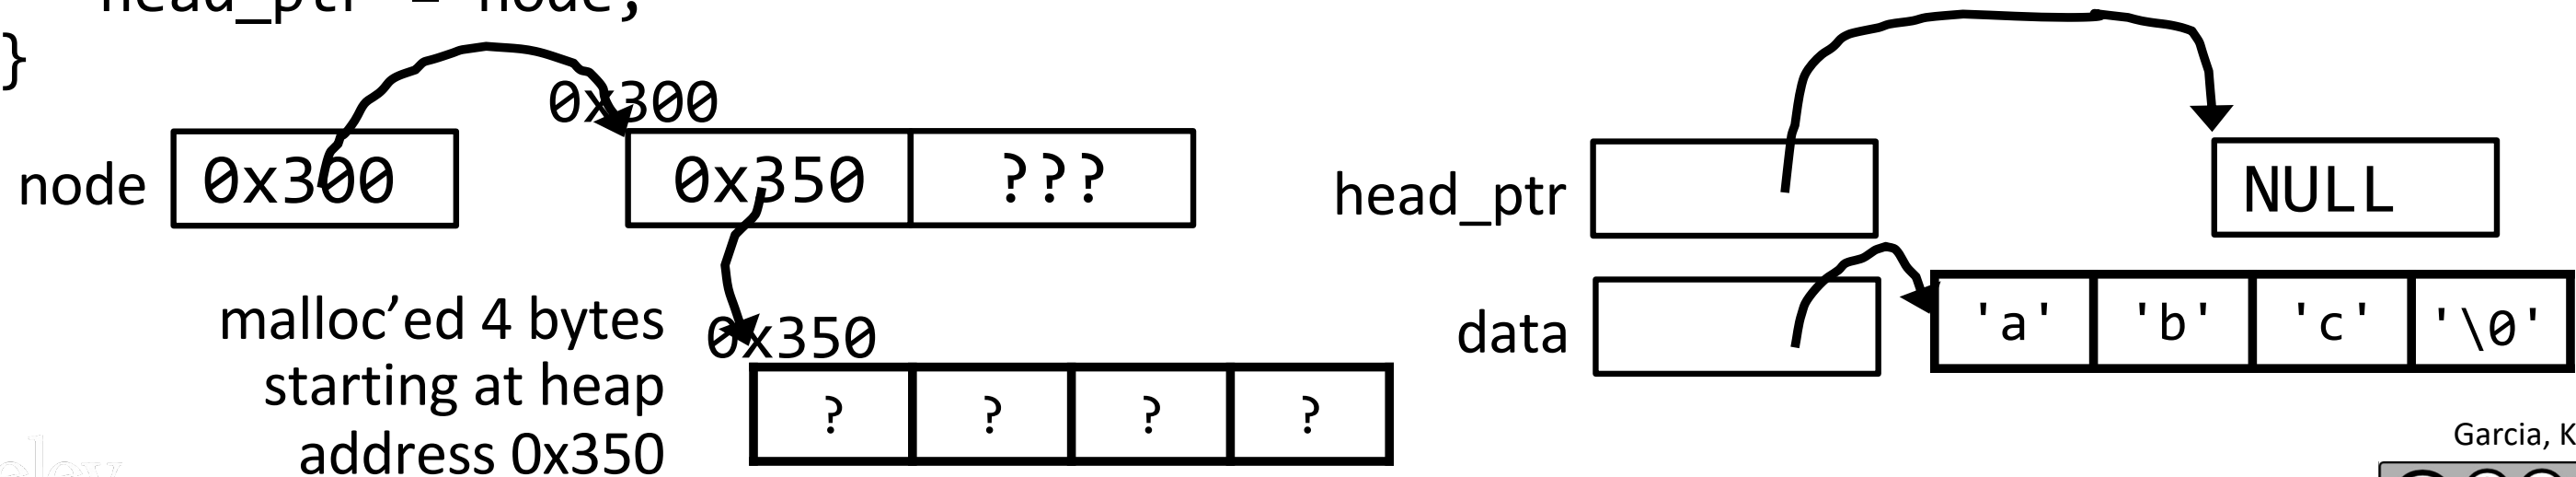



Linked List Example

```
# include <string.h>
int main() {
1   struct Node *head = NULL;
2   add_to_front(&head, "abc");
   ... // free nodes, strings here...
}

void add_to_front(struct Node **head_ptr, char *data) {
3   struct Node *node = (struct Node *) malloc(sizeof(struct Node));
4   node->data = (char *) malloc(strlen(data) + 1); // extra byte!
5   strcpy(node->data, data); // strcpy also copies null terminator
6   node->next = *head_ptr;
7   *head_ptr = node;
}
```

```
struct Node {
    char *data;
    struct Node *next;
};
```



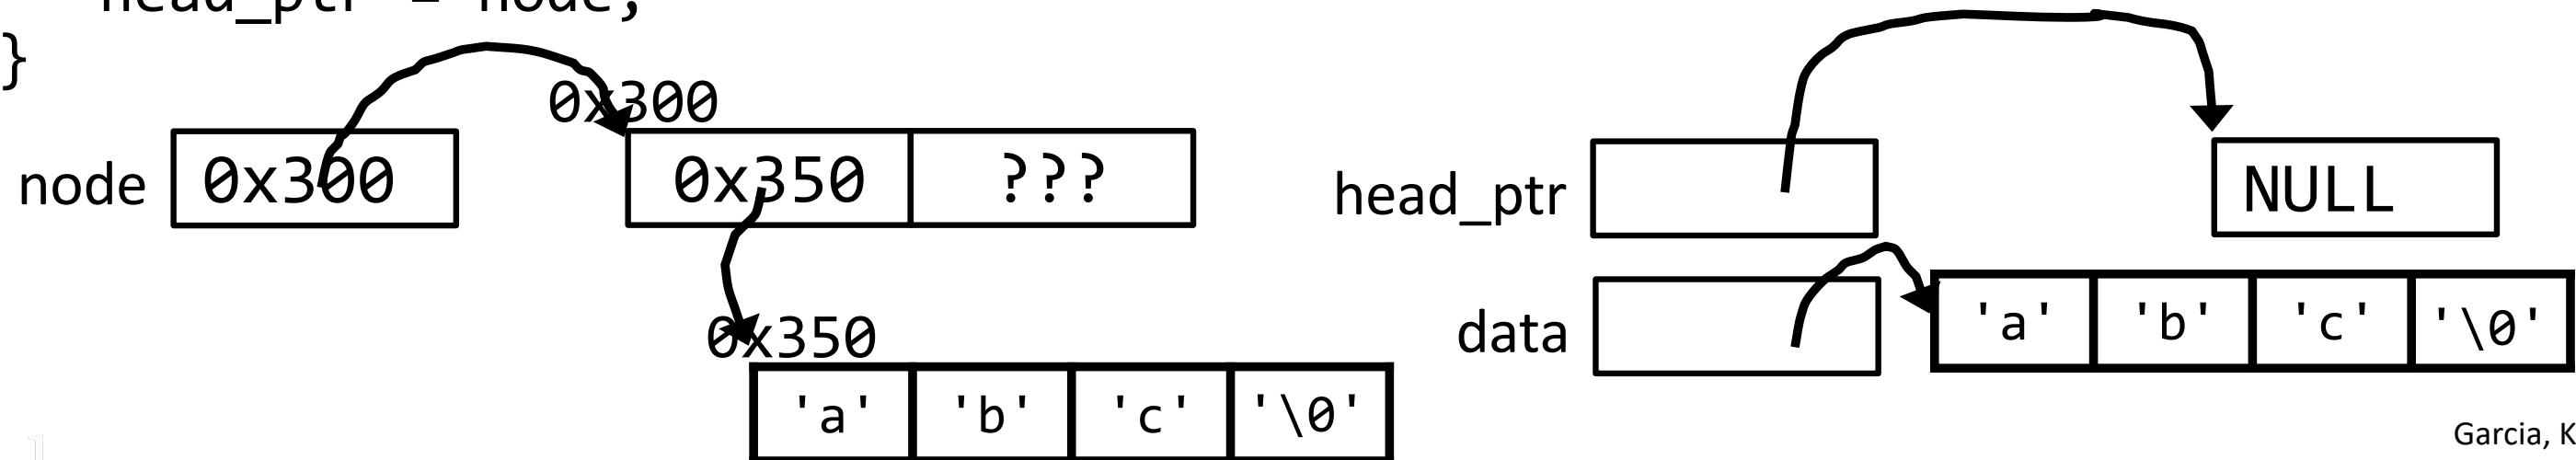



Linked List Example

```
# include <string.h>
int main() {
1   struct Node *head = NULL;
2   add_to_front(&head, "abc");
   ... // free nodes, strings here...
}

void add_to_front(struct Node **head_ptr, char *data) {
3   struct Node *node = (struct Node *) malloc(sizeof(struct Node));
4   node->data = (char *) malloc(strlen(data) + 1); // extra byte!
5   strcpy(node->data, data); // strcpy also copies null terminator
6   node->next = *head_ptr;
7   *head_ptr = node;
}
```

```
struct Node {
    char *data;
    struct Node *next;
};
```



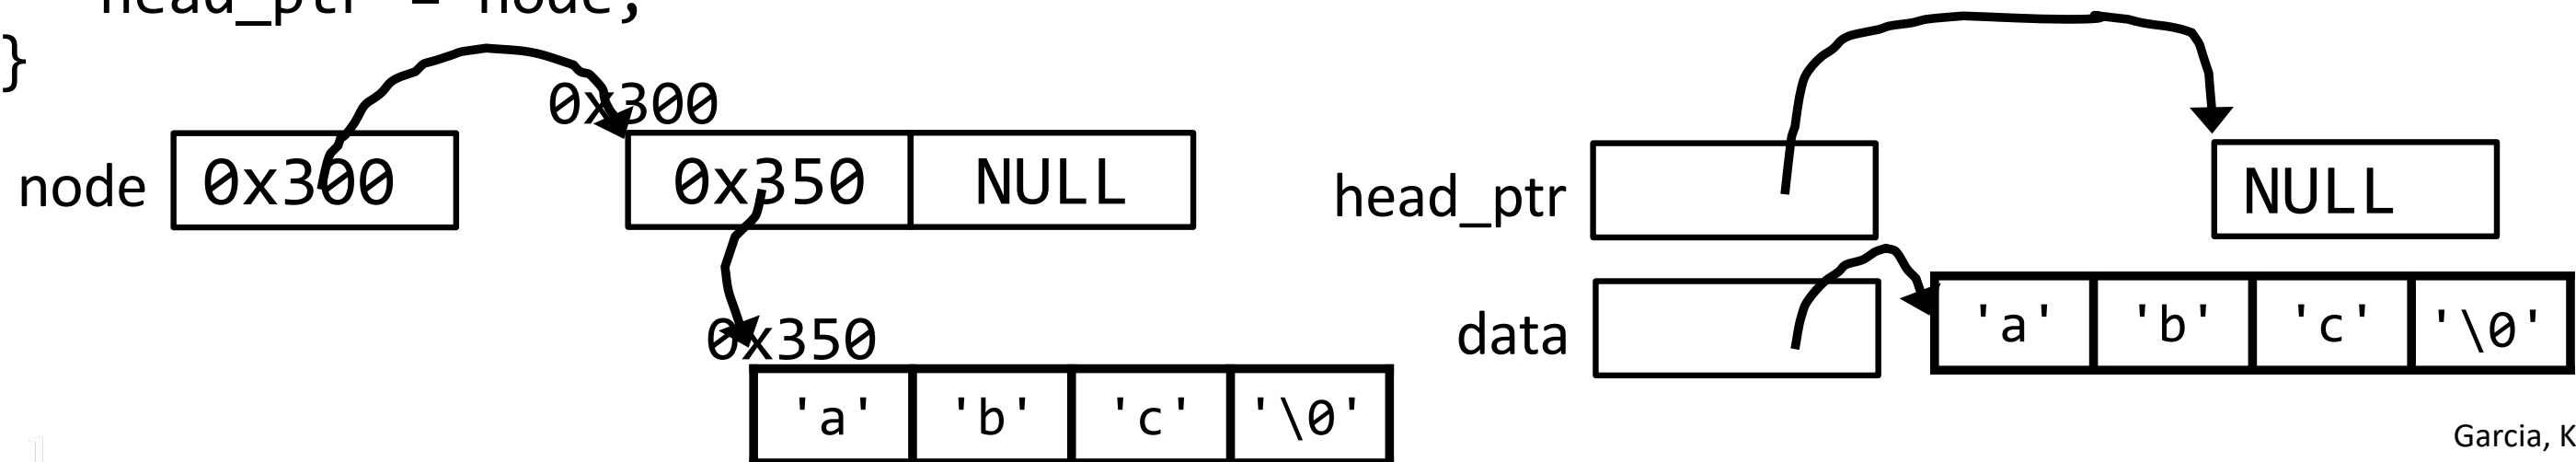



Linked List Example

```
# include <string.h>
int main() {
1   struct Node *head = NULL;
2   add_to_front(&head, "abc");
   ... // free nodes, strings here...
}

void add_to_front(struct Node **head_ptr, char *data) {
3   struct Node *node = (struct Node *) malloc(sizeof(struct Node));
4   node->data = (char *) malloc(strlen(data) + 1); // extra byte!
5   strcpy(node->data, data); // strcpy also copies null terminator
6   node->next = *head_ptr;
7   *head_ptr = node;
}
```

```
struct Node {
    char *data;
    struct Node *next;
};
```



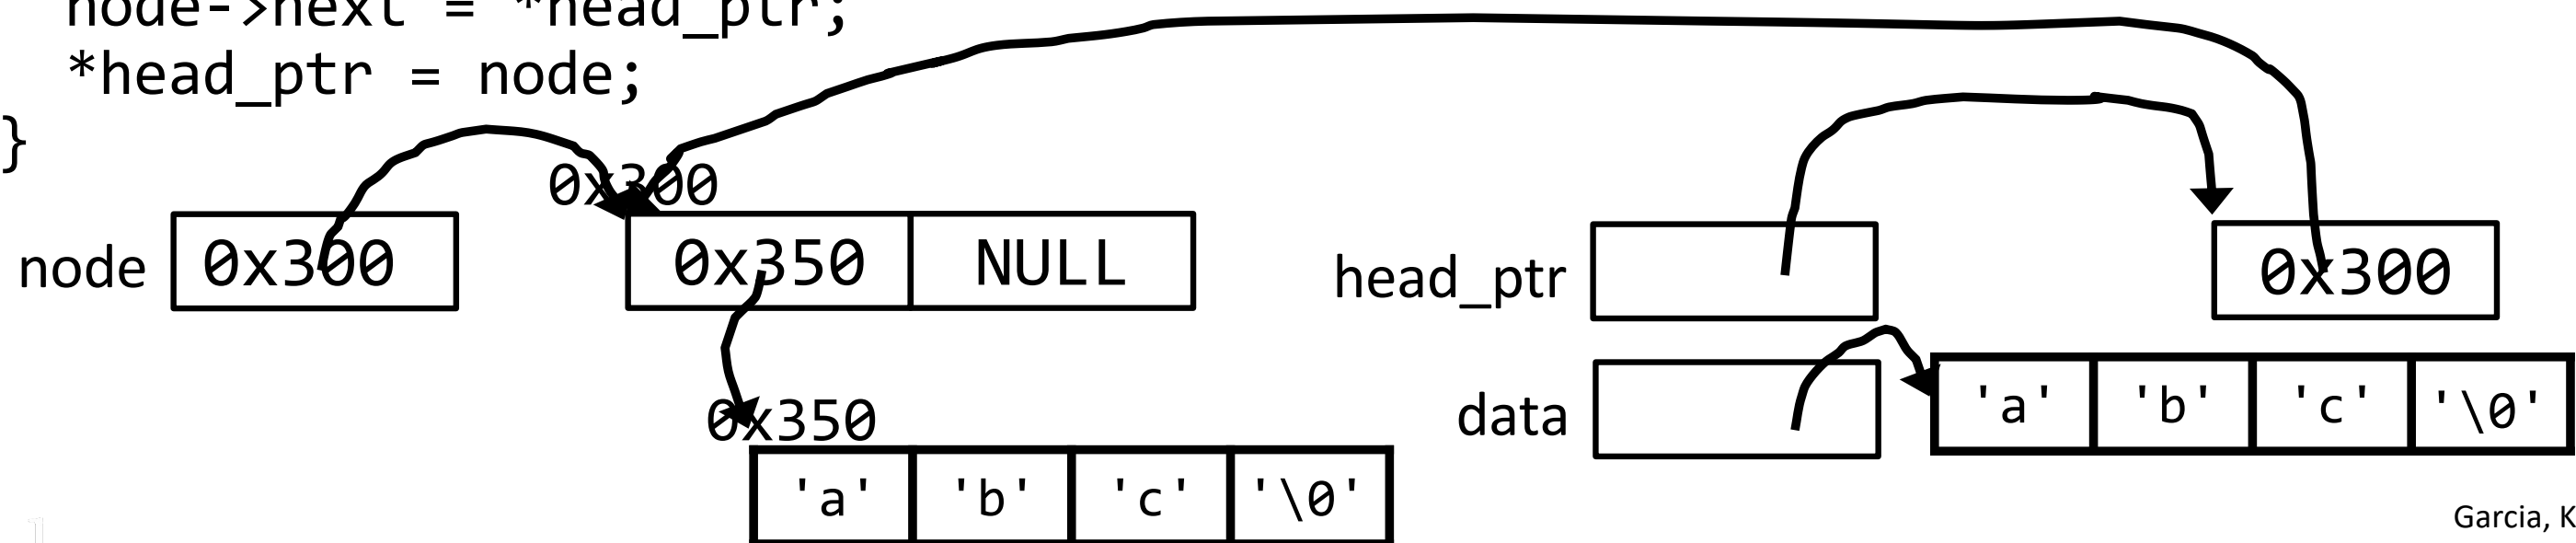



Linked List Example

```
# include <string.h>
int main() {
1   struct Node *head = NULL;
2   add_to_front(&head, "abc");
   ... // free nodes, strings here...
}

void add_to_front(struct Node **head_ptr, char *data) {
3   struct Node *node = (struct Node *) malloc(sizeof(struct Node));
4   node->data = (char *) malloc(strlen(data) + 1); // extra byte!
5   strcpy(node->data, data); // strcpy also copies null terminator
6   node->next = *head_ptr;
7   *head_ptr = node;
}
```

```
struct Node {
    char *data;
    struct Node *next;
};
```






Linked List Example


```
# include <string.h>
int main() {
1   struct Node *head = NULL;
2   add_to_front(&head, "abc");
   ... // free nodes, strings here...
}

void add_to_front(struct Node **head_ptr, char *data) {
3   struct Node *node = (struct Node *) malloc(sizeof(struct Node));
4   node->data = (char *) malloc(strlen(data) + 1); // extra byte!
5   strcpy(node->data, data); // strcpy also copies null terminator
6   node->next = *head_ptr;
7   *head_ptr = node;
}
```

```
struct Node {
    char *data;
    struct Node *next;
};
```

Check out the
lecture code in
Drive!






Agenda

When Memory Goes Bad

- Memory Locations
- The Stack
- The Heap
- Linked List Example
- When Memory Goes Bad
- Implementing Memory Management



Working with Memory

Code, Static storage are easy:

- They never grow or shrink.

Stack space is also easy:


- Stack frames are created and destroyed in LIFO order.
- Just avoid “**dangling references**”: pointers to deallocated variables (e.g., from old stack frames).

⚠ Working with the heap is tricky:


- Memory can be allocated / deallocated at any time!
- “**Memory leak**”: If you forget to deallocate memory } Your program will eventually run out of memory

- “**Use after free**”: If you use data after calling free
 - “**Double free**”: If you call free 2x on same memory
- } Possible crash or exploitable vulnerability

0xFFFF FFFF




0x0000 0000



Failure to free()

- The runtime does not check for the programmer's failure to manage memory.
 - Memory is so performance-critical that there just isn't time to do this.
 - Usual result: you corrupt the memory allocator's internal structure, and you find out much later in a totally unrelated part of your code!
- **Memory leak: Failure to free() allocated memory**
 - 🙄 Initial symptoms. Nothing...
 - Until you hit a critical point, memory leaks aren't actually a problem
 - 📈 ...Later symptoms: performance drops off a cliff...
 - Memory hierarchy behavior tends to be great just up until it isn't, then it hits several cliffs
 - 😈 ...and then your program is killed off!
 - Because the operating system (OS) says "no" when you ask for more memory



Use after Free




- **“Dangling reference”**

When you keep using a pointer, even after it has been deallocated

```
struct foo *f;  
...  
f = malloc(sizeof(struct foo));  
...  
free(f);  
...  
bar(f->a); // !!!
```

- Reads after the free may be corrupted!
 - If something else takes over that memory, your program will probably read the wrong information!
- Writes corrupt other data!
 - Uh oh... Your program crashes later!




Double-Free...



```
struct foo *f = (struct foo *) malloc(10*sizeof(struct foo));  
...  
free(f);  
...  
free(f); // !!!
```

- May cause either a use-after-free (because something else called `malloc()` and got that data) or corrupt heap data (because you are no longer freeing a pointer tracked by `malloc`)



Forgetting realloc() Can Move Data

- "Dangling reference"
- Remember, when you realloc it can copy data to a different part of the heap.


```
int *nums;
nums = malloc(10*sizeof(int));
...
int *g = nums;
...
nums = realloc(nums,
               20*sizeof(int));

// g could now point
// to invalid memory
```

```
int *nums;
nums = malloc(10*sizeof(int));
...
// forget to update nums
// on realloc call
realloc(nums, 20*sizeof(int));

// nums could now point
// to invalid memory,
// and we could have potentially
// lost a pointer to a new block
```

- Reads may be corrupted, and writes may corrupt other pieces of memory.



Faulty Heap Management

[Concept Check]

How many memory management errors are in this code?


[for next time]

```
void free_mem_x() {  
    int fnh[3];  
    ...  
    free(fnh);  
}  
  
void free_mem_y() {  
    int *fum = malloc(4*sizeof(int));  
    free(fum+1);  
    ...  
    free(fum);  
    ...  
    free(fum);  
}
```

- A. 1
- B. 2
- C. 3
- D. 0
- E. Other

How many memory management errors are in this code?





Faulty Heap Management

[Concept Check]

How many memory management errors are in this code?

[for next time]

```
void free_mem_x() {  
    int fnh[3];  
    ...  
    free(fnh);  
}
```


} free() on stack-allocated memory

```
void free_mem_y() {  
    int *fum = malloc(4*sizeof(int));  
    free(fum+1);  
    ...  
    free(fum);  
    ...  
    free(fum);  
}
```

} free() on memory that isn't the pointer from malloc

} Double free()


- A. 1
- B. 2
- C. 3
- D. 0
- E. Other



Valgrind to the rescue!!!

- Valgrind slows down your program by an order of magnitude, but...
 - It adds a tons of checks designed to catch most (but not all) memory errors
 - Memory leaks
 - Misuse of free
 - Writing over the end of arrays
- Tools like Valgrind are absolutely essential for debugging C code.

Check out Lab 02!




And in Conclusion...

- C has 3 pools of memory for variables:
 - Data: **global/static** variable storage, basically permanent
 - Stack: **local variable** storage, parameters, return address
 - Heap (**dynamic** storage): malloc() grabs space from here, free() returns it.
 - (4th memory pool: text, for the program executable itself)
- Heap data is biggest source of bugs in C code!



Garcia, Kao




Agenda

Implementing Memory Management

- Memory Locations
- The Stack
- The Heap
- Linked List Example
- When Memory Goes Bad
- Implementing Memory Management


Material not tested. Recording:
<https://www.youtube.com/watch?v=Sq5tSeWfnGY>



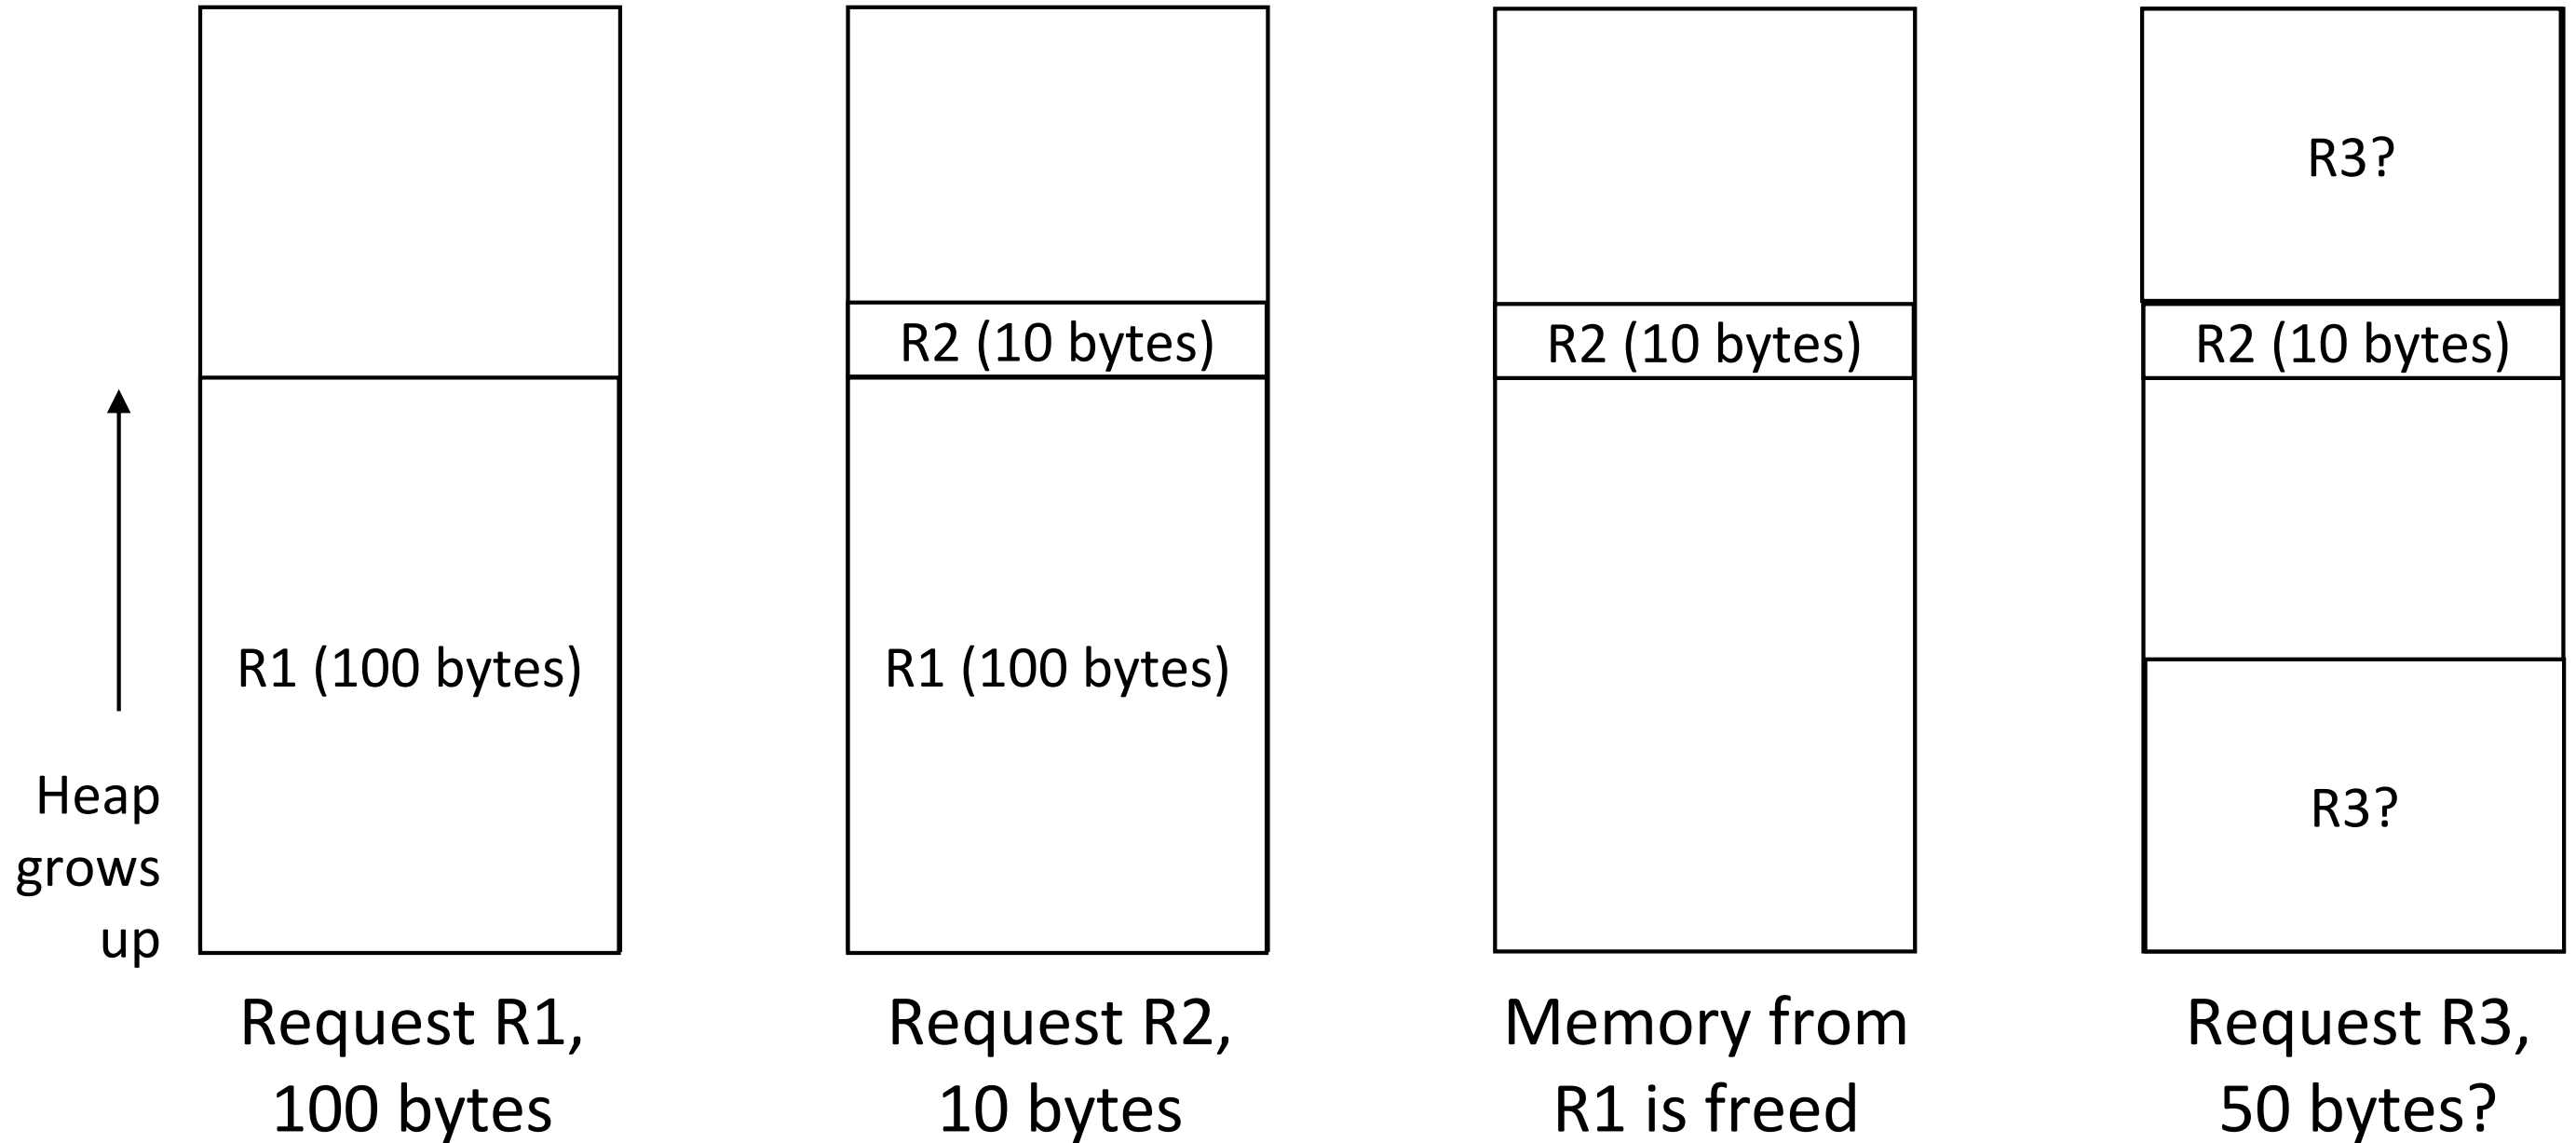
Heap Management Requirements


- Want `malloc()` and `free()` to run quickly
- Want minimal memory overhead
- Want to avoid fragmentation*,
when most of our free memory is in many small chunks
 - In this case, we might have many free bytes but not be able to satisfy a large request since the free bytes are not contiguous in memory.

* This is technically
external fragmentation




Heap Management Example






K&R Malloc/Free Implementation

- From Section 8.7 of K&R
 - Code in the book uses some C language features we haven't discussed and is written in a very terse style; don't worry if you can't decipher the code
- Each block of memory is preceded by a header that has two fields:
 - size of the block, and
 - a pointer to the next block
- All free blocks are kept in a circular linked list.
- In an allocated block, the header's pointer field is unused.



K&R Malloc/Free Implementation

- `malloc()` searches the free list for a block that is big enough. If none is found, more memory is requested from the operating system. If what it gets can't satisfy the request, it fails.
- `free()` checks if the blocks adjacent to the freed block are also free.
 - If so, adjacent free blocks are merged (**coalesced**) into a single, larger free block.
 - Otherwise, freed block is just added to the free list.



Choosing a block in `malloc()`

- If there are multiple free blocks of memory that are big enough for some request, how do we choose which one to use?
 - **best-fit**: choose the smallest block that is big enough for the request.
 - **first-fit**: choose the first block we see that is big enough.
 - **next-fit**: like first-fit, but remember where we finished searching and resume searching from there.